

Instruction Selection & Scheduling via SMT

CS517

Raine Wheary, Christine Lin
[GitHub project](#)

1) Introduction

We consider the problem of backend code generation in a compiler. By this point in the pipeline, we can understand the input program as having been reduced to a graph of basic operations in some intermediate representation (IR). The backend must further lower these operations into machine instructions.

For linear control flow (i.e. within a single basic block), there are three major tasks to be solved:

- **Instruction selection:** Modern CPU architectures often have complex macroinstructions encoding more than one basic operation. For example, the x86-64 instruction `mov rax, [rcx + 4*rdx + 28]` involves a left shift, two additions, and a load. Other examples include stack operations, multiple loads and stores (e.g. ARM's `ldp` and `stp`), and fused-multiply-add. A compiler will prefer to emit these instructions when applicable, but it can be a net loss if it means intermediate results are no longer available and need to be recomputed. In the last example, if the address value `rcx + 4*rdx + 28` was used in many places, a compiler should compute this effective address with `lea` and reuse it, rather than doing the calculation inside the `mov`.
- **Instruction scheduling:** Even when two machine instructions could be independent of each other, their order relative to each other can impact the performance of the generated code. We model a pipelined but not out-of-order processor, where a pipeline stall occurs if an instruction's inputs are not available.
- **Register allocation:** Programs in IR form use an unlimited number of virtual registers, but these need to be mapped to a limited number of physical registers on the machine.

Each of these tasks should be understood as a search problem. Classically, the backend is split into three phases, with instruction selection, instruction scheduling, and register allocation being performed in order [1]. However, this architecture has its limitations. When deciding whether to emit a higher-latency macroinstruction, there are situations where the optimal choice depends on register pressure or on the ability of the scheduler to avoid a pipeline stall.

Register allocation is well known to be NP-complete, and is typically solved with heuristic algorithms in practice. We will focus on an integrated algorithm for first two problems via reduction to SMT.

2) Problem statement

We will adopt the convention of coloring words and variables related to the IR (the input) **blue** and those related to the machine program (the output) **red**.

2.1) The IR program

We assume there is a fixed finite set of **IR opcodes** C and each opcode $\alpha \in C$ has an associated **arity**.

As input, we start with an **IR program**, which is a sequence P of **IR instructions**. Each **IR instruction** consists of an opcode and a tuple of arguments matching the opcode's arity.

$$P := \{P_1, \dots, P_N\}$$
$$P_i ::= \alpha(t_1, \dots, t_m), \quad \alpha \in C, \quad m = \text{arity}(\alpha), \quad 1 \leq t_1, \dots, t_m < i.$$

The requirement $1 \leq t_1, \dots, t_m < i$ captures the idea that the input program is a DAG, because each parameter of an **IR instruction** is the index of some previous instruction. This also means that the first instruction must be nullary.

In addition to the sequence of instructions, a subset of program indices $R \subseteq \{1, \dots, N\}$ are designated as **results**. An index $i \in R$ being present in the set means that the result of the instruction P_i must be materialized in the **machine program**, i.e. it cannot be coalesced into an intermediate computation of a **machine instruction**.

For each instruction P_i in a given **IR program**, we can build an **expression tree** (denoted $\text{tree}(P_i)$) by recursively replacing the indices t in the parameter list with the nodes they refer to:

$$\begin{aligned} \text{tree}(P_i) &:= \alpha(\text{tree}(P_{t_1}), \dots, \text{tree}(P_{t_m})) \\ \text{where } P_i &= \alpha(t_1, \dots, t_m) \end{aligned}$$

(Expression trees cannot actually be materialized in an implementation, since they may be exponentially large without sharing, but they are useful concept in the specification of the problem.)

2.2) The **machine program**

We have a finite set C' of **machine opcodes** $\beta \in C'$ with associated arities. A **machine program** M has a similar structure to the input program: it consists of a sequence $\{M_1, \dots, M_K\}$ of **machine instructions**, each of which consists of an opcode and a number of previous argument references that matches the arity.

$$\begin{aligned} M &:= \{M_1, \dots, M_K\} \\ M_j &:= \beta(r_1, \dots, r_m), \quad \beta \in C', \quad m = \text{arity}(\beta), \quad 1 \leq r_1, \dots, r_m < j. \end{aligned}$$

As a convention, we use variables i and t as indexes of **IR instructions** and j and r for **machine instructions**.

A **machine instruction** is thought of as standing in for one or more **IR instructions**. For example, a 3-ary machine opcode **FMA**(a, b, c) might represent the computation **add**($a, \text{mul}(b, c)$). We formalize this with the idea of a **machine definition** D , which associates each n -ary machine opcode β with a definition in terms of a tree of **IR instructions** with n free variables, written τ .

$$\begin{aligned} \tau &::= x_k && \text{(free variable)} \\ &| \alpha(\tau_1, \dots, \tau_m) && \alpha \in C, \quad m = \text{arity}(\alpha) \\ \text{fv}(\tau) &:= \begin{cases} \{x_k\} & \text{if } \tau = x_k \\ \bigcup_{k=1}^m \text{fv}(\tau_k) & \text{if } \tau = \alpha(\tau_1, \dots, \tau_m) \end{cases} && \text{(set of free variables of a tree } \tau) \\ D_\beta &::= ((x_1, \dots, x_m), \tau) && \beta \in C', \quad m = \text{arity}(\beta), \quad \text{fv}(\tau) = \{x_1, \dots, x_m\} \end{aligned}$$

Given a machine definition $D_\beta = ((x_1, \dots, x_m), \tau)$, we can apply it to a concrete m -tuple of trees, written as $D_\beta(\tau_1, \dots, \tau_m)$, by substituting each x_k for τ_k in the body of the definition τ .

Given a **machine program** and a collection of machine definitions, we can define an analogous notion of **expression tree** for each instruction M_j .

$$\begin{aligned} \text{tree}(D, M_j) &:= D_\beta(\text{tree}(D, M_{r_1}), \dots, \text{tree}(D, M_{r_m})) \\ \text{where } M_j &= \beta(r_1, \dots, r_m) \end{aligned}$$

For a fixed C, C', D , we say that an instruction M_j in a **machine program** **models** an **IR instruction** P_i if $\text{tree}(P_i) = \text{tree}(D, M_j)$. In other words, M_j computes the same expression tree as P_i .

Then we can say that a **machine program** M **models** an **IR program** (P, R) if for each designated result $i \in R$ there exists some instruction M_j that models P_i . In other words, every **IR instruction** designated as a result is computed by some **machine instruction**. Also note that the modeling requirement does not forbid the **machine program** from containing unnecessary instructions.

Notice that if an index i is not present in R then there is no requirement that the **machine program** materialize the value of P_i . This is what allows the compiler to use instructions like **FMA**(a, b, c) when we don't need the intermediate computation **mul**(b, c) to ultimately be saved in a register.

2.3) Latency calculation

Each machine opcode $\beta \in C'$ has an associated **latency**, which we take to be a positive integer. To model an in-order pipelined CPU, each instruction $M_j = \beta(r_1, \dots, r_m)$ has the following life cycle:

- **Decode:** The instruction is fetched and decoded by the CPU. This happens the cycle after the previous instruction M_{j-1} is dispatched.
- **Dispatch:** This happens as soon as the instruction is decoded and all the dependencies $M_{r_1}, \dots, M_{r_m}$ have been retired.
- **Retire:** This happens L cycles after the instruction is dispatched, where $L = \text{latency}(\beta)$.

A **machine program** is considered finished once all instructions have been retired; this determines its **total latency**. In notation:

$$\begin{aligned}
 \text{decode}(j) &:= \begin{cases} 0 & \text{if } j = 1 \\ \text{dispatch}(j-1) + 1 & \text{if } j > 1 \end{cases} \\
 \text{dispatch}(j) &:= \max\{\text{decode}(j), \text{retire}(r_1), \dots, \text{retire}(r_m)\} \\
 &\quad \text{where } M_j = \beta(r_1, \dots, r_m) \\
 \text{retire}(j) &:= \text{dispatch}(j) + \text{latency}(\beta) \\
 &\quad \text{where } M_j = \beta(r_1, \dots, r_m) \\
 \text{latency}(M) &:= \max\{\text{retire}(j) \mid 1 \leq j \leq K\} \\
 &\quad \text{where } M = \{M_1, \dots, M_K\}
 \end{aligned}$$

2.4) Putting it together

We now have all the pieces to define instruction selection and scheduling as a search problem:

Problem. Given

- A set of IR opcodes C and their arities,
- An IR program P and set of results R ,
- A set of machine opcodes C' and their arities,
- A map D_β associating opcodes $\beta \in C'$ with machine definitions,
- A map $\text{latency}(\beta)$ associating opcodes $\beta \in C'$ with their latency,

Find a **machine program** M such that M models (P, R) and $\text{latency}(M)$ is minimized.

There is a corresponding decision problem: given the above and a latency bound L , determine whether there is a **machine program** M such that $\text{latency}(M) \leq L$ is NP-hard. We claim that the decision problem is NP-hard, even if the values of $\text{latency}(\beta)$ and L are given in unary.

In our implementation, only P , D , and $\text{latency}(-)$ need to be provided by the user. The sets C and C' and associated arities are automatically deduced, and R is hardcoded to be $\{N\}$, i.e. only the last IR instruction is designated as a result. This is not a significant restriction, because if a larger set $R = \{i_1, \dots, i_n\}$ is desired, the user can introduce a new n -ary IR opcode **tuple** and add the “dummy” instruction **tuple** $(i_1, \dots, i_n)$ at the end of the program, along with a corresponding machine opcode **TUPLE** with $D_{\text{TUPLE}}(x_1, \dots, x_n) = \text{tuple}(x_1, \dots, x_n)$. Since there is no way to compute **tuple** $(\dots)$ except using **TUPLE** $(\dots)$, this forces $P_{i_1}, \dots, P_{i_n}$ to be materialized.

3) Design of the reduction to SMT

References

- [1] G. S. H. Blindell, “Survey on Instruction Selection: An Extensive and Modern Literature Review.” [Online]. Available: <https://arxiv.org/abs/1306.4898>